





**OPEN SOURCE SUMMIT** 

**China 2023** 

# SIG Instrumentation: Intro, Deep Dive and Recent Developments

Mengjiao Liu Shivanshu Raj Shrivastava

## About us





Mengjiao Liu@DaoCloud

SIG Instrumentation Reviewer
SIG Docs Reviewer & Chinese localization Owner

Github: mengjiao-liu Twitter: Mengjiao99



Shivanshu Raj Shrivastava

Contributor: K8S, OpenTelemetry, Istio

Twitter: shivanshu1333 Github: shivanshu1333

# Agenda



- Introduction
- SIG Subprojects
- Recent Developments
  - Metrics
  - Logs
  - Traces
  - Demo
- Get involved

### What is SIG Instrumentation?



Defined in <u>SIG Instrumentation Charter</u>: Owns best practices for cluster observability through metrics, logging, events, and traces across all k8s components & development of components required for all k8s clusters.

### Scope:

- Logs
- Events
- Metrics
- Traces

### SubProjects:

- kube-state-metrics
- klog
- metrics-server
- o and more!

### How do we do it?



- Triage and fix relevant instrumentation issues
  - All open SIG Instrumentation-labelled issues and pull requests
- Review all code changes for metrics, logs, events and traces
- Develop new features and enhancements
  - Kubernetes Enhancement Proposals (KEPs) for SIG Instrumentation
- Maintain and support subprojects
- Mentor new contributors



\_\_\_\_\_ China 2023

# Subprojects

# Subprojects



- klog
- kube-state-metrics
- metrics-server
- prometheus-adapter
- usage-metrics-collector





Package <u>klog</u> implements logging analogous to the Google-internal C++ INFO/ERROR/V setup. It provides functions Info, Warning, Error, Fatal, plus formatting variants such as Infof. It also provides V-style logging controlled by the -v and -vmodule=file=2 flags.

The <u>kubernetes/kubernetes</u> repo uses klog as the logging library

- → Supports structured logging (include text and json format)
- → Supports contextual logging

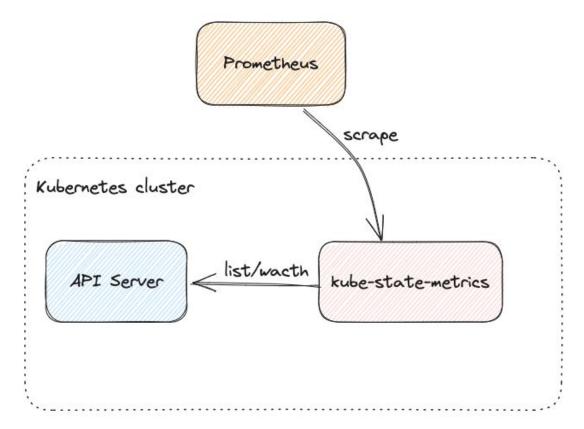
**& Release**: <u>v2.100.1</u>

## kube-state-metrics



<u>kube-state-metrics</u> generates Prometheus format metrics based on the current state of the Kubernetes native resources

 Deployments, Pods, Services, StatefulSets etc. A full list of resources is available in the documentation of kube-state-metrics.



### **Examples:**

kube\_deployment\_status\_condition{namesp
ace="kube-system",deployment="coredns",c
ondition="Available",status="true"} 1

kube\_pod\_status\_ready{namespace="defaul
t",pod="nginx",uid="2d3795ca-e53d-4bd3-86
Of-ced3e74a3db9",condition="true"} 1



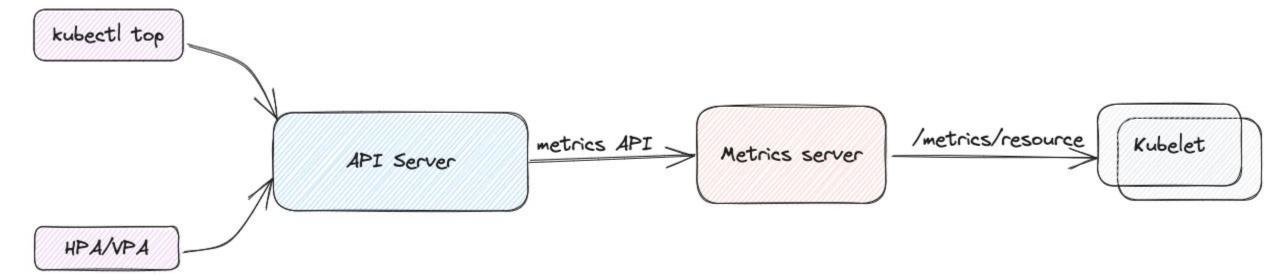
Release: <u>v2.10.0</u>

### metrics-server



Metrics Server is a scalable, efficient source of container resource metrics for Kubernetes built-in autoscaling pipelines.

- offers a basic set of metrics to support automatic scaling and similar use cases
- view these metrics using the kubectl top command



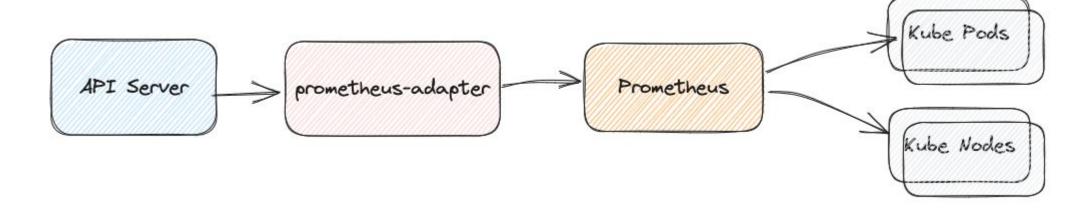
🎊 Release: v0.6.4

# prometheus-adapter



<u>Prometheus Adapter</u> Implementation of the Kubernetes Custom, Resource and External Metric APIs

- suitable for use with the autoscaling/v2 Horizontal Pod Autoscaler
- replace the metrics server on clusters that already run Prometheus



🎉 Release: <u>v0.11.1</u>

# usage-metrics-collector



- Repo: <a href="https://github.com/kubernetes-sigs/usage-metrics-collector">https://github.com/kubernetes-sigs/usage-metrics-collector</a>
- Prometheus metrics collector optimized for collecting kube usage and capacity metrics.
  - Simple
    - Don't require writing complex promptl statements to combine metrics –
       native support for joining metadata across resources
  - Scalable
    - Support large clusters with lots of Pods
    - Optimize cardinality and storage requirements by exporting pre-aggregated results
  - Performant
    - Quickly compute values without requiring tiering prometheus instances

# Example Collector Configurations



Get p95 utilization (cpu and memory) using 1 second sampling intervals for all containers in each workload.

```
resources:
  "cpu": "cpu cores" # get cpu metrics
  "memory": "memory bytes" # get memory metrics
aggregations:
- sources:
    type: "container"
    container: [ "utilization" ] # export container utilization
  levels:
  - mask:
      name: "container"
     builtIn: # aggregate on these labels
       exported container: true
       exported namespace: true
       workload name: true
       workload kind: true
       workload api group: true
       workload api version: true
      operation: "p95" # take the 95th percentile sample
```

### **Result Metrics:**

```
workload_p95_utilization_cpu_cores{exported_container="", exported_namespace="", workload_name="",
workload_kind="", workload_api_group="", workload_api_version=""}

workload_p95_utilization_memory_bytes{exported_container="", exported_namespace="", workload_name="",
workload_kind="", workload_api_group="", workload_api_version=""}
```



Recent Developments





OPEN SOURCE SUMMIT

**China 2023** 

# Metrics



### KEP-3498: Extend metrics stability



Enhancement status

Alpha: 1.26

Beta: 1.27

Stable: 1.28

- Extend on the existing metrics stability framework to achieve stability guarantees
- Document all metrics exported by the core distribution of Kubernetes

### **KEP-3498**: Extend metrics stability



### **Stability Levels:**

- Internal: does not have any stability guarantees
- Alpha: does not have any stability guarantees
  - included in metric <u>auto-documentation</u>
- Beta: has some stability guarantees
  - deprecate metrics need a release or more
  - forward compatible, labels can only be added, and not removed
  - included in metric <u>auto-documentation</u>
- **Stable**: has stability guarantees
  - follow a deprecation policy(12 months or 3 releases)
  - not change in respect to labels
  - included in metric <u>auto-documentation</u>



Documentation Kubernetes Blog Training Partners Community Case Studies Versions - English -

Q Search

- Documentation
- Getting started
- Concepts
- ▶ Tasks
- Tutorials
- Reference

#### Glossary

- API Overview
- ▶ API Access Control
- Well-Known Labels, Annotations and Taints
- Kubernetes API
- Instrumentation

Service Level Indicator Metrics

CRI Pod & Container Metrics

Node metrics data

#### **Kubernetes Metrics Reference**

- Kubernetes Issues and Security
- ▶ Node Reference Information
- Networking Reference
- Setup tools
- ▶ Command line tool (kubectl)
- Component tools
- Debug cluster
- Configuration APIs
- External APIs
- Scheduling
- Other Tools
- Contribute

Kubernetes Documentation / Reference / Instrumentation / Kubernetes Metrics Reference

#### **Kubernetes Metrics Reference**

#### Metrics (v1.29)

This page details the metrics that different Kubernetes components export. You can query the metrics endpoint for these components using an HTTP scrape, and fetch the current metrics data in Prometheus format.

#### List of Stable Kubernetes Metrics

Stable metrics observe strict API contracts and no labels can be added or removed from stable metrics during their lifetime.

#### apiserver\_admission\_controller\_admission\_duration\_seconds

Admission controller latency histogram in seconds, identified by name and broken out for each operation and API resource and type (validate or admit).

- Stability Level: STABLE
- Type: Histogram
- · Labels: name operation rejected type

#### apiserver\_admission\_step\_admission\_duration\_seconds

Admission sub-step latency histogram in seconds, broken out for each operation and API resource and step type (validate or admit).

- Stability Level: STABLE
- · Type: Histogram
- · Labels: operation rejected type

#### apiserver\_admission\_webhook\_admission\_duration\_seconds

Admission webhook latency histogram in seconds, identified by name and broken out for each operation and API resource and type (validate or admit).

- Stability Level: STABLE
- Type: Histogram
- Labels: name operation rejected type

• Create an issue (auto-generated page)

Print entire section

Metrics (v1.29)

List of Stable Kubernetes Metrics List of Beta Kubernetes Metrics

List of Alpha Kubernetes Metrics







S OPEN SOURCE SUMMIT

China 2023

Logs





- Structured Logging
  - Future of observability in Kubernetes
  - Goals/Non Goals
  - Implementation details
  - Migration details



- GA in v1.27; introduced as Alpha in v1.19
- Migration
  - Most components have been migrated.
  - The remaining part is to convert code directly to contextual logging.



### Goals

- Make most common logs more queryable by standardizing log message and references to Kubernetes objects (Pods, Nodes etc.)
- Enforce log structure by introduction of new klog methods that could be used to generate structured logs.
- Simplify ingestion of logs into third party logging solutions by adding an option to output logs in the JSON format

### Non Goals

 We are not replacing currently used logging library (klog) or the way in which it is used



- Implementation details
  - Log message structure
  - References to Kubernetes objects
  - Introduce JSON output format in klog
  - Logging configuration
  - Performance
- Migration details



- Log message structure
  - There could be multiple ways of standardising the logging structure in Kubernetes, the one we agreed to implement in the KEP is to have following logging structure

```
<message> <key1>=<value1> <key2>=<value2> ...
e.g.
pod := corev1.Pod{Name: "kubedns", Namespace: "kube-system", ...}
klog.InfoS("Pod status updated", "pod", klog.KObj(pod), "status",
"ready")
```



- References to Kubernetes objects:
  - The idea is to use k8s api first approach to get k8s objects and embed the object related information into the logs
  - Correlate between different kubernetes objects

```
func KObj(obj ObjectMeta) ObjectRef func KRef(namespace, name string) ObjectRef
```

```
type ObjectRef struct {
  Name     string `json:"name"`
  Namespace string `json:"namespace,omitempty"`
}
```



References to Kubernetes objects:

```
Namespaced objects: <namespace>/<name>
e.g. kube-system/kubedns
Non-namespaced objects: <name>
e.g. node cluster1-vm-72x33b8p-34jz
e.g.
klog.InfoS("Pod status updated", "pod", klog.KObj(pod), "status",
"ready")
```

klog.ErrorS(err, "Failed to update pod status", "pod", klog.KObj(pod))



- Introduce JSON output format in klog:
  - Introduction of new methods to klog library to support JSON.
  - With klog v2 we can take further advantage of this fact and add an option to produce structured logs in JSON format.
- Some pros of using JSON:
  - Broadly adopted by logging libraries with very efficient implementations (zap, zerolog).
  - Out of the box support by many logging backends (Elasticsearch, Stackdriver, BigQuery, Splunk, Open Telemetry)
  - Easily parsable and transformable
  - Existing tools for ad-hoc analysis (jq)



Introduce JSON output format in klog:

```
klog.InfoS("Pod status updated", "pod", klog.KObj(pod), "status",
"ready")
 "ts": 1580306777.04728,
 "v": 4,
 "msg": "Pod status updated",
 "pod":{
    "name": "nginx-1",
    "namespace": "default"
 "status": "ready"
```



Introduce JSON output format in klog:

```
type Request struct {
 Method string
 Timeout int
 secret string
 Con *Connection
req := Request{Method: "GET", Timeout: 30, secret: "pony"}
klog.InfoS("Request finished", "request", Request)
```



Introduce JSON output format in klog:

```
"ts": 1580306777.04728,
"v": 4,
"msg": "Request finished",
"request":{
 "Method": "GET",
 "Timeout": 30
```



- Logging configuration
  - Implementation of LoggingConfig structure as part of k8s.io/component-base

introduced flag --logging-format values:

- a) text: for text-based logging format (default)
- b) json: for new JSON format



### Performance

 Logging performance with the new implementation. Performance is wrt to log volume and the performance impact.

logger	time [ns/op]	bytes[B/op]	allocations[alloc/op]
Text Infof	2252	248	3
Text InfoS	2455	280	3
JSON Infof	1406	19	1
JSON InfoS	319	67	1

- InfoS implementation for text is 9% slower than Infof.
- Kubernetes performance as logging takes less than 2% of overall CPU usage.



- Beta in v1.28; introduced as Alpha in v1.24
- Migration
  - kube-controller-manager has been converted
  - kube-scheduler has converted most of it
  - Other components are being converted



- Contextual Logging
  - Goals/Non Goals
  - Implementation details
  - Migration details



- Contextual Logging
  - Goals
    - Grant the caller of a function control over logging inside that function, either by passing a logger into the function or by configuring the object that a method belongs to.
    - Provide documentation and helper code for setting up logging in unit tests.
    - Change as few exported APIs as possible
  - Non-Goals
    - Remove the klog text output format
    - Deprecate klog



- Removing the dependency on the global klog logger
  - klog.ErrorS -> logger.Error
  - o logger is a logr.Logger instance. klog.Logger is an alias for that type
- Extend klog for contextual logging
  - Several new klog functions help with that:
    - klog.FromContext
    - klog.Background
    - klog.TODO

### Contextual Logging



```
// FromContext retrieves a logger set by the caller or, if not set,
// falls back to the program's global logger (a Logger instance or klog
// itself).

func FromContext(ctx context.Context) Logger {
        if logging.contextualLoggingEnabled {
            if logger, err := logr.FromContext(ctx); err == nil {
                return logger
            }
        }
        return Background()
}
```

### Contextual Logging



- Extend klog for contextual logging
  - Use migrated structured logs and attach context with it
    - With the helper functions that we added in klog, structured logs can also be transformed to more fine grained contextual logs







S OPEN SOURCE SUMMIT

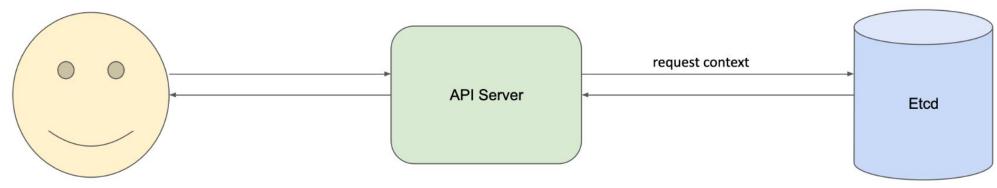
China 2023

# Traces



### Tracing in Kubernetes









### API Server Tracing:

- Beta in 1.27
- Trace requests from the API Server to Etcd.

### **Kubelet Tracing**

- Beta in 1.27
- Trace requests from the Kubelet to the Container

Runtime

OpenTelemetry dependency updated to 1.0+ in K8s 1.26



### Log-Based "Tracing"

Trace[1395114870]: "Create" url:/api/v1/nodes,user-agent:tracing.test/v0.0.0 (linux/amd64)

kubernetes/\$Format,audit-id:c0282104-8068-44b1-a088-756b1253326d,client:127.0.0.1,accept:application/vnd.kubernetes.protob

uf, \*/\*,protocol:HTTP/2.0

(28-Oct-2022 13:42:35.876) (total time: 2ms):

Trace[1395114870]: ---"limitedReadBody succeeded" len:86 0ms (13:42:35.876)

Trace[1395114870]: ---"About to convert to expected version" 0ms (13:42:35.876)

Trace[1395114870]: ---"Conversion done" 0ms (13:42:35.876)

Trace[1395114870]: ---"About to store object in database" 0ms (13:42:35.876)

Trace[1395114870]: ["Create etcd3"

audit-id:c0282104-8068-44b1-a088-756b1253326d,key:/minions/fake,type:\*core.Node,resource:nodes 2ms (13:42:35.876)

Trace[1395114870]: ---"About to Encode" 0ms (13:42:35.876)

Trace[1395114870]: ---"Encode succeeded" len:177 0ms (13:42:35.876)

Trace[1395114870]: ---"TransformToStorage succeeded" 0ms (13:42:35.876)

Trace[1395114870]: ---"Txn call succeeded" 1ms (13:42:35.878)

Trace[1395114870]: ---"decode succeeded" len:177 0ms (13:42:35.878)]

Trace[1395114870]: ---"Write to database call succeeded" len:86 0ms (13:42:35.878)

Trace[1395114870]: ---"About to write a response" 0ms (13:42:35.878)

Trace[1395114870]: ---"Writing http response done" 0ms (13:42:35.878)

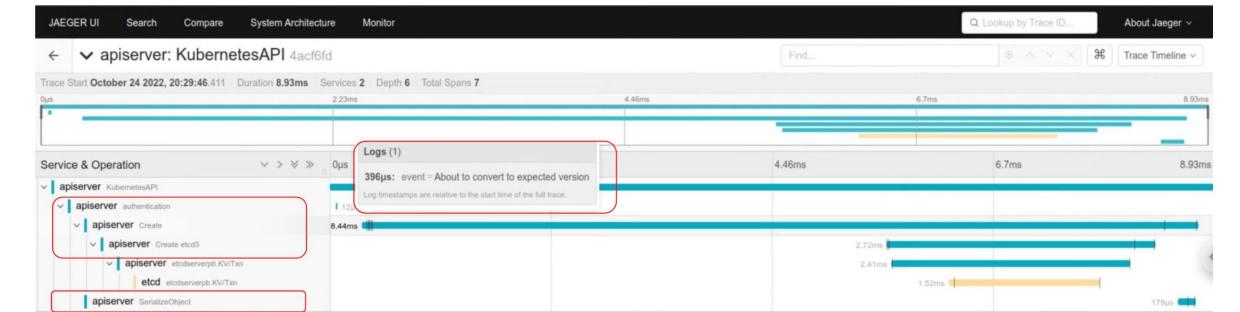
Trace[1395114870]: [2.771143ms] [2.771143ms] END



#### Kubernetes 1.22

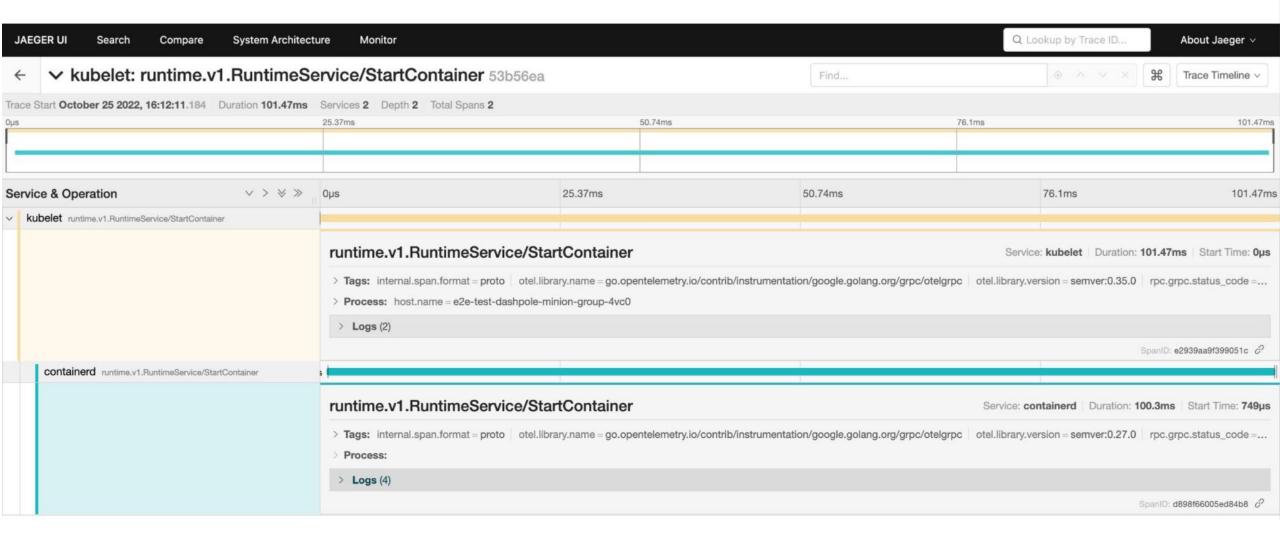


#### Kubernetes 1.26



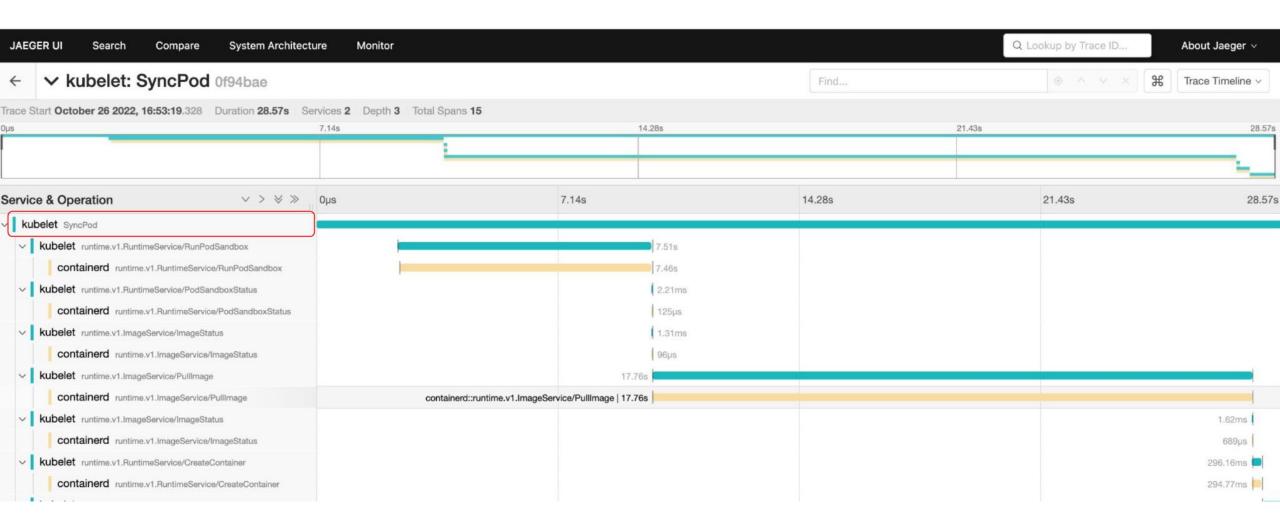


### Alpha: CRI Traces





### **Proof of Concept: Complete Pod traces**





### **Future Plans:**

 Add kubelet spans to track "create pod" instead of just "create

### container"

- Link from Metrics to Traces with Prometheus Exemplars
- Link from Logs to Traces with Trace + Span IDs in Logs







S OPEN SOURCE SUMMIT

China 2023

# Demo



Get involved!

**China 2023** 

### Get involved



- Join <u>#sig-instrumentation</u> on Slack
- Attend our SIG meetings
  - Regular SIG Meeting: Thursdays at 9:30 PT (Pacific Time)
     (biweekly).
  - Regular Triage Meeting: Thursdays at 9:30 PT (Pacific Time) (biweekly - alternating with regular meeting)
- Subscribe mailing list: <u>kubernetes-sig-instrumentation</u>
- Participate in <u>reviews</u>, <u>issues</u>, and <u>docs</u>!
- Kubernetes Enhancement Proposals (KEPs)

We are seeking more contributors! You are welcome to join us!

