



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2023

How Multi-Cloud Multi-Cluster HPA Helps Trip.com Group Deal with Business Downturn and Rapid Recovery

Honghui Yue&Jingxue Li

Agenda

- Background
- Design
 - Reliability First
 - Capacity Scheduling
- Implementation
- Lessons learned

About Us

Trip.com Group

Trip.com Group is a leading global travel service provider comprising of Trip.com, Ctrip, Skyscanner, and Qunar with the mission "to pursue the perfect trip for a better world".

Across its platforms, Trip.com Group helps travellers around the world make informed and cost-effective bookings for travel products and services. Trip.com Group technology enables partners to connect their offerings with users through the aggregation of comprehensive travel-related content and resources, and an advanced transaction platform consisting of apps, websites and 24/7 customer service centers.



Project Core Contributors



Honghui Yue
Container&Hybrid Cloud
Cloud Native R&D Director & Architect



Jingxue Li
Container&Hybrid Cloud
Cloud Native R&D Expert

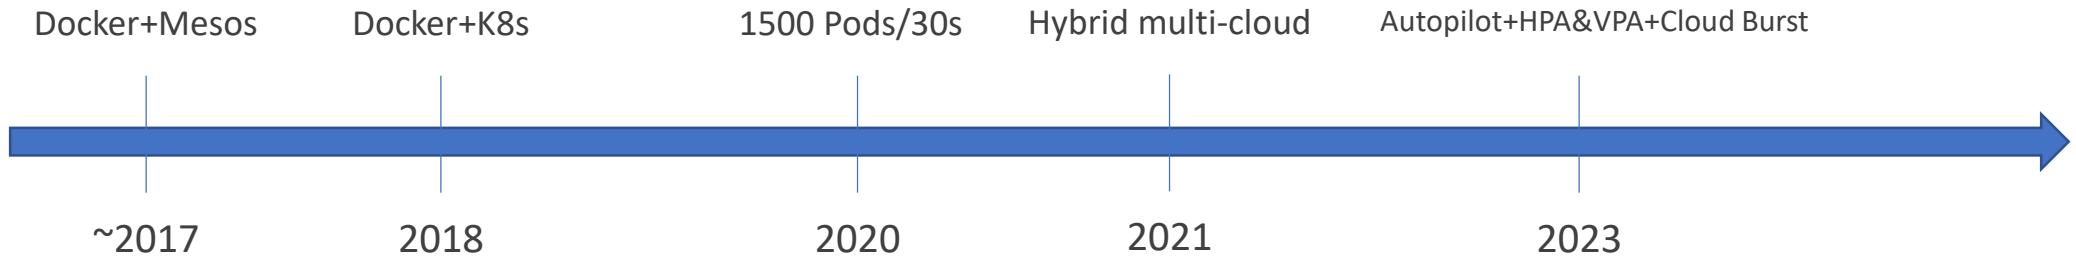


Zongqing Li
Container&Hybrid Cloud
Senior Cloud Native R&D Engineer



Yao Wang
Container&Hybrid Cloud
Cloud Native R&D Engineer

The Cloud Native Trip



Kubernetes Native

~30 K8s Clusters
~500K Pods
~20 Workloads(App、Redis、Kafka...)

Hybrid Multicloud

~5 Cloud Providers
~3 Regions(Asia, EU)
~60 Cloud Native Services

Elastic Scheduling

1500 Pods/30s
90%+ Apps enable
Autopilot+HPA&VPA
App Node avg CPU 25%~35%

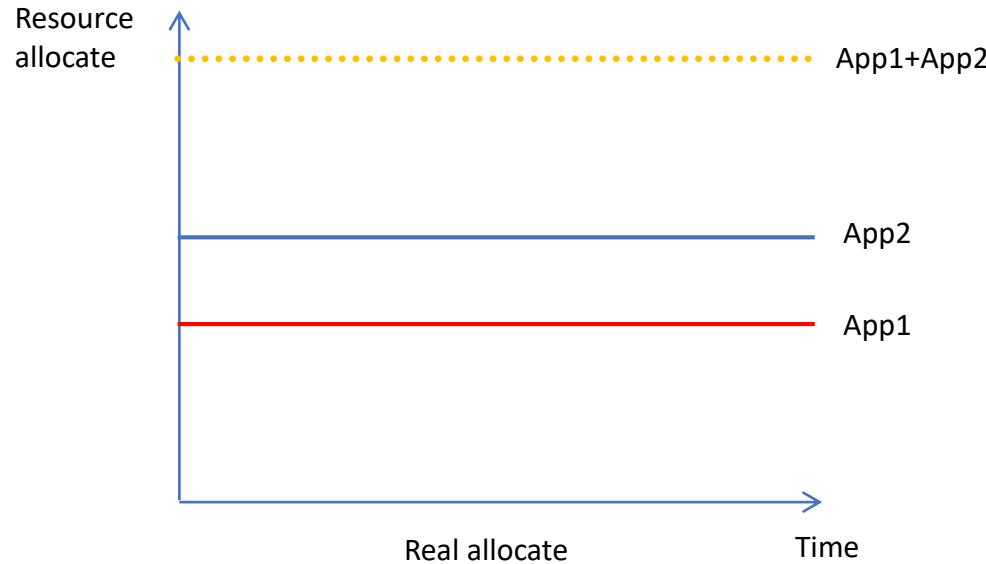
The Cloud Native Trip

- Decreased the capacity peak of AZ by more than **28%** during the last Labor Day
- We successfully dealt with multiple holiday traffic peaks after the pandemic, with the highest expansion peak of **~2000 Pod/min**
- More than **30%** of load was offloaded to the public cloud at peak times

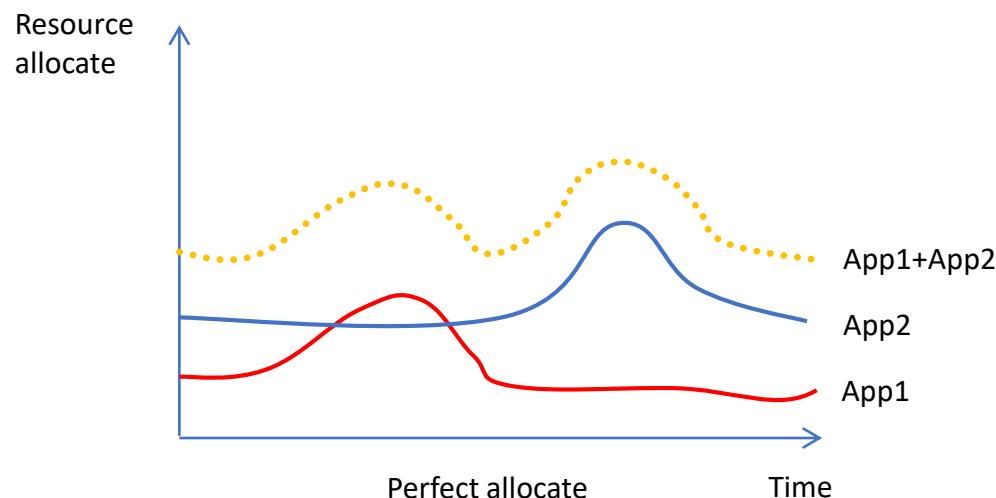
What problems can HPA solve?

- It's hard to maintain the capacity of Apps manually
 - Why
 - Too many deployments: **10k+** apps and **20k+** deployments
 - Hard to predict traffic for individual Apps: Covid-19 and policy changing
 - Hard to balance between availability and cost: high pressure on cost saving during pandemic, which means less resource redundancy.
 - Bad result
 - **Low efficiency:** Too much time is wasted on maintaining capacity of Apps.
 - **Capacity risk:** Disordered and paniced scales of Apps during the site peak put an overwhelming strain on resources.
 - **Bad Experience:** Developers&SREs are unhappy.

What problems can HPA solve?



- Too much resources are wasted
 - Why
 - Scale out beyond the actual demand.
 - Scale out in advance, without resource time-sharing effect between Apps.
 - Bad result
 - Low resource utilization(avg CPU < 10%)
 - Need to prepare much more resources than the actual demand for the site capacity plan.



Agenda

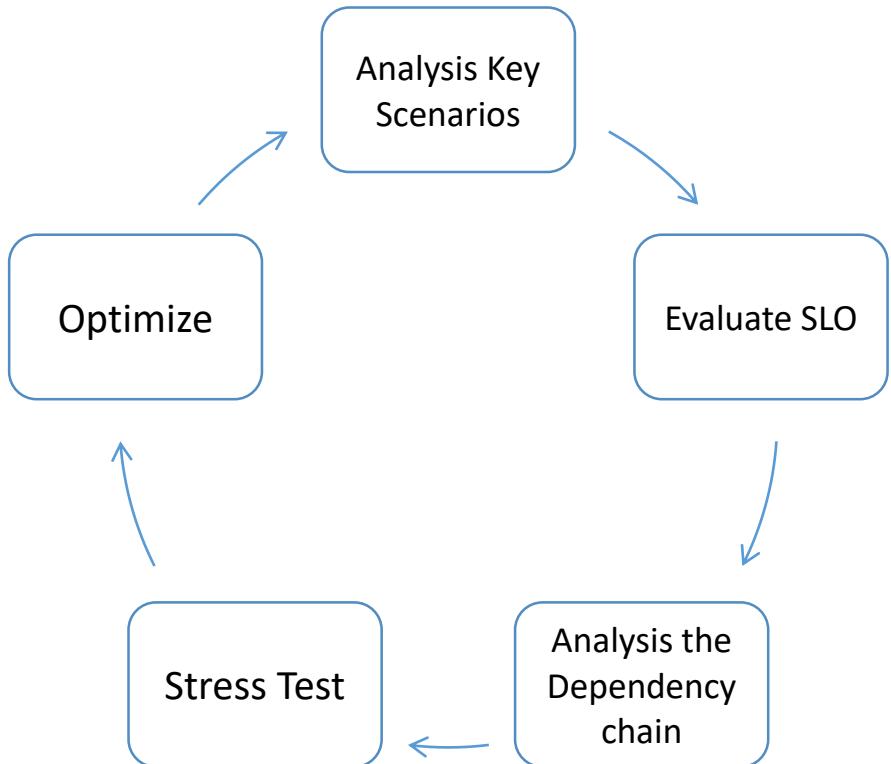
- Background
- Design
 - Reliability First
 - Capacity Scheduling
- Implementation
- Lessons learned

Confidence&Trust

HPA looks good, but what if it fails to scale at critical moments?

High performance + high reliability + enough capacity!

E2E Scale Performance SLO



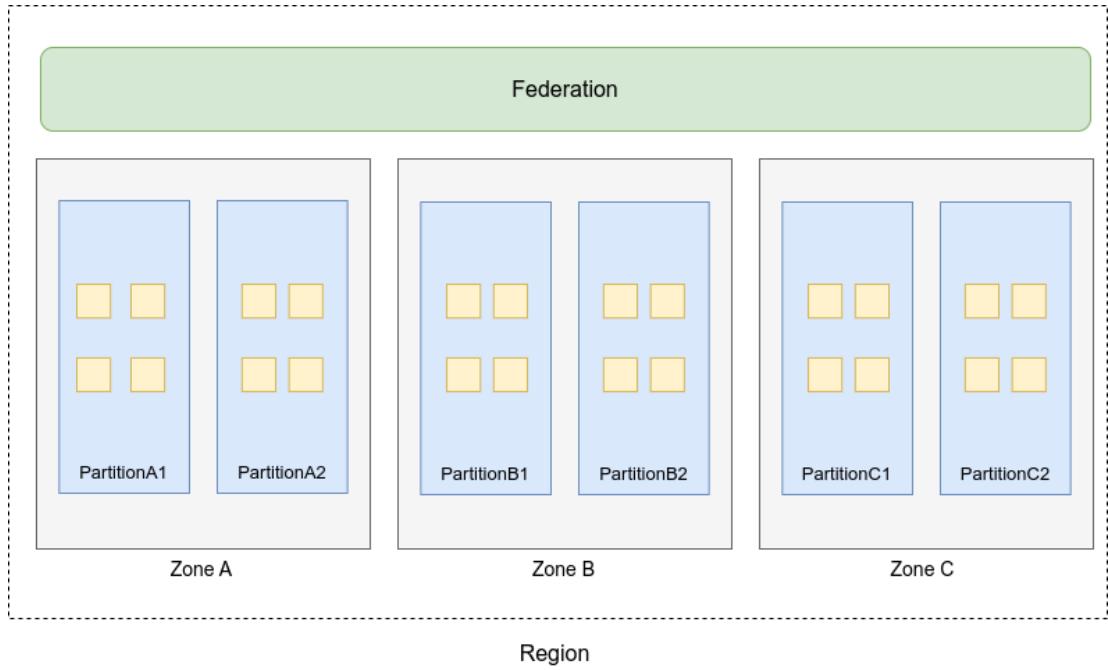
- Evaluate SLO based on scenarios
 - Daily Expansion
 - Holidays
 - Disaster Recovery
 - Future growth
- Stress testing
 - Construct testcases based on product data
 - Test e2e in real environment
 - Collect key metrics(throughput, latency, success rate)
- Optimization
 - Refactor the architecture and minimize cross-cluster dependency in the east-west direction.
 - Optimize the code

Fault isolation architecture

How can we ensure the reliability of the overall system, as partial failure is unavoidable?

Fault isolation+Redundancy+Failover!

Fault isolation architecture



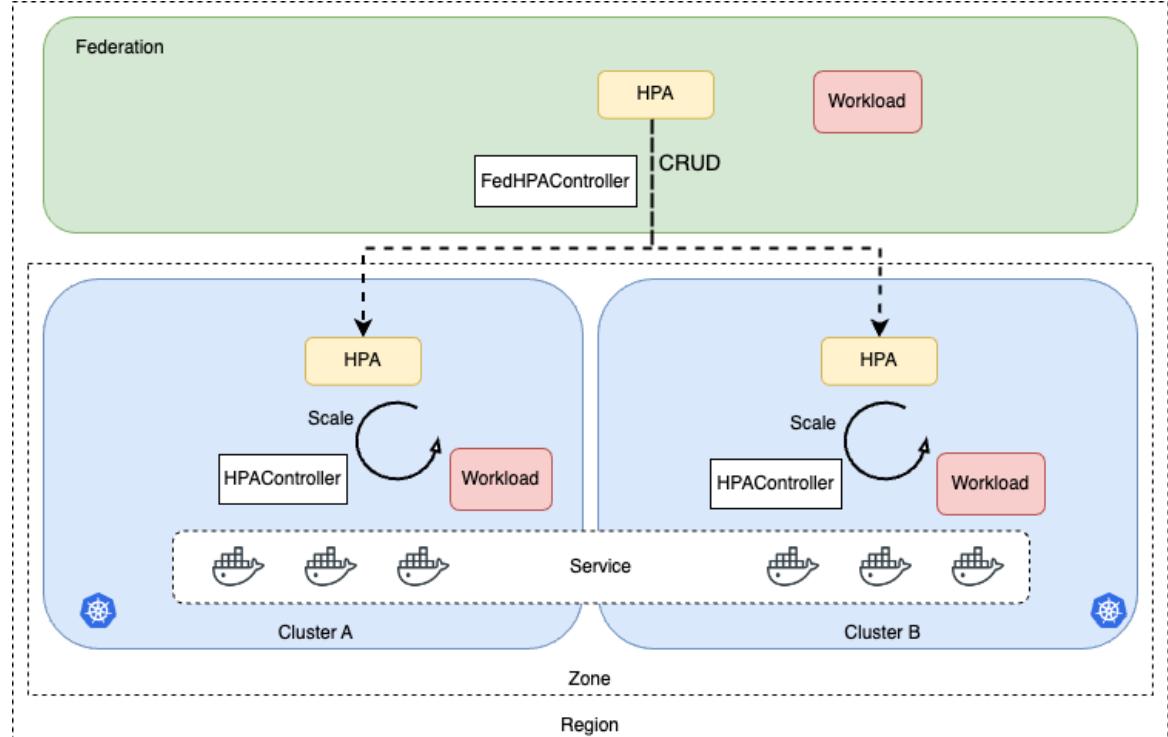
- **Architecture**

- Isolation domain: Region/AZ/Partition
- At least 1 K8s cluster in each partition, and 2 partitions in each AZ, at least 3 AZs in each region
- The regional Federation is responsible for
 - Hiding the K8s member clusters to the end users
 - Managing AZ/Partition/Cluster within the region
 - Scheduling resource between Clusters within the AZ

- **Key constraints**

- The critical processing chain should be done within the fault domain
- Minimize dependency between isolation domains in the east-west direction
- App pods should be distributed across multiple partitions within multiple AZs
- Changes should not be released on different AZs or even different partitions at the same time

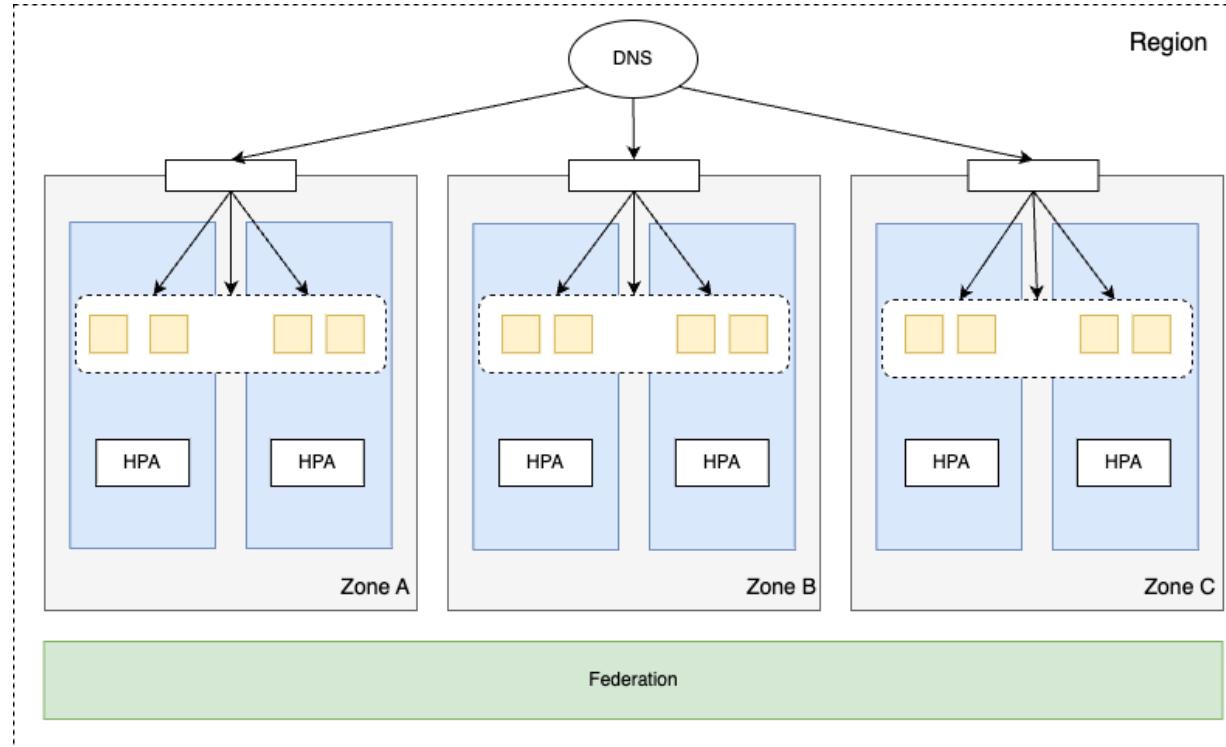
Fault isolation architecture



Isolation domains for multi-cluster HPA

function	critical	frequency	isolation domain
workload scale	yes	high	Cluster
HPA CR CRUD	no	low	Region
cluster rebalance	no	low	Region

Capacity scheduling on Hybrid-cloud



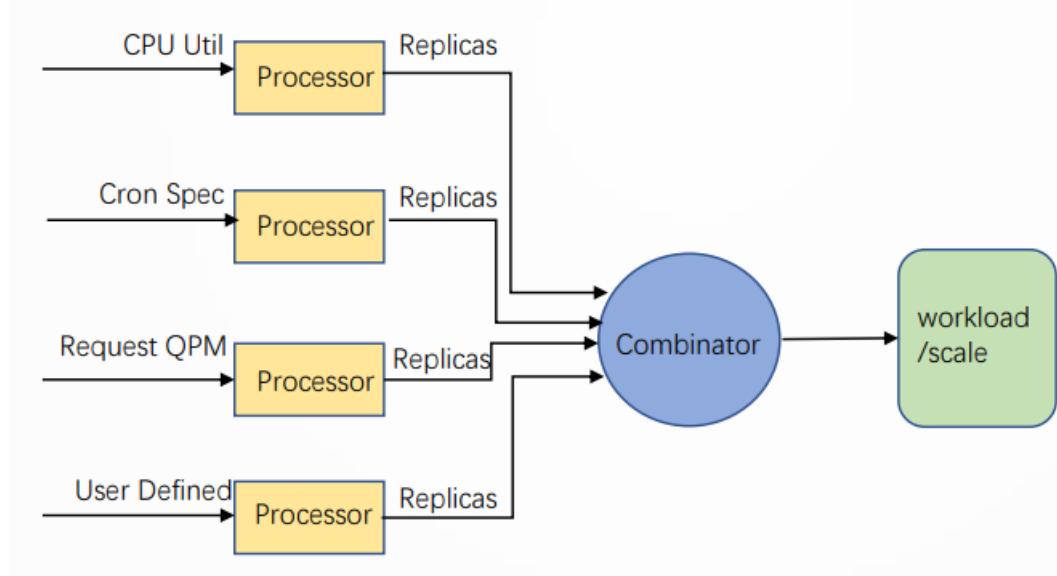
Zone A: Public Cloud
Zone B&C: Private Cloud

- Network&Traffic
 - Routable pod IP across clusters
 - App pods within the same AZ share the same traffic scheduling unit(e.g.: svc, ingress)
 - Locality load balancing
- Capacity scheduling
 - Schedule load between AZs by traffic load balancing
 - Rebalance between clusters within each AZ by Federation scheduling

Agenda

- Background
- Design
 - Reliability First
 - Capacity Scheduling
- Implementation
- Lessons learned

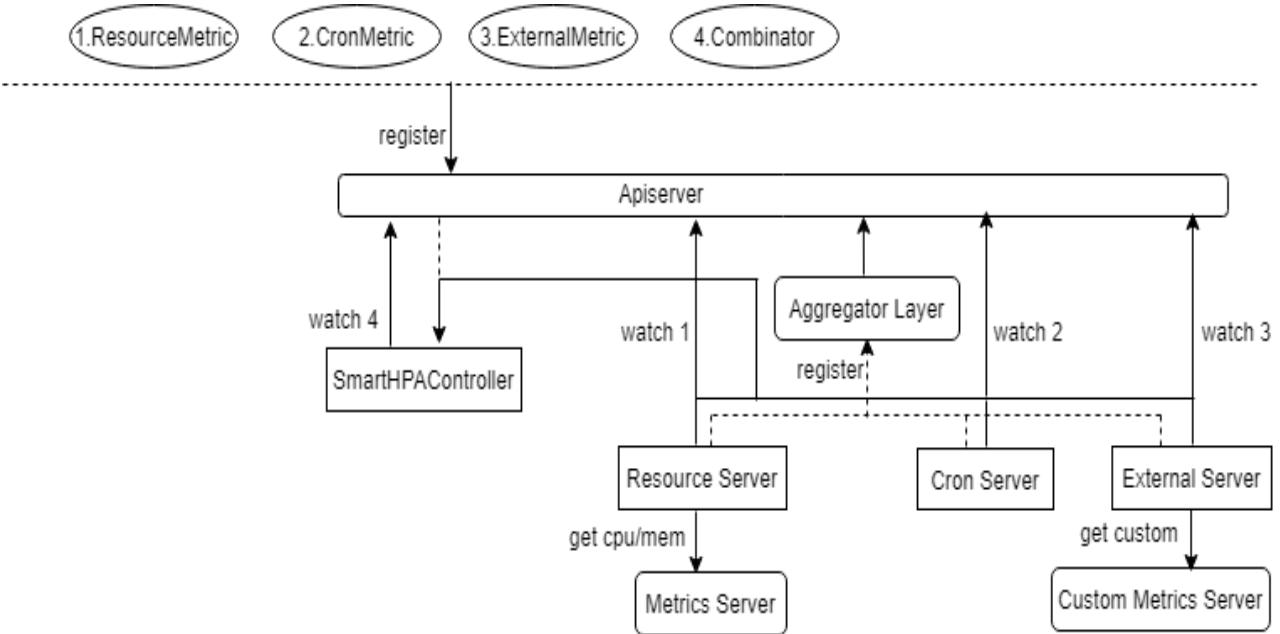
Design of HPA



- Components

- Processor
 - output replicas recommendation
- Combinator
 - calculate the final replicas leverage the input of replicas recommendations from Processor
 - Scale the target workload

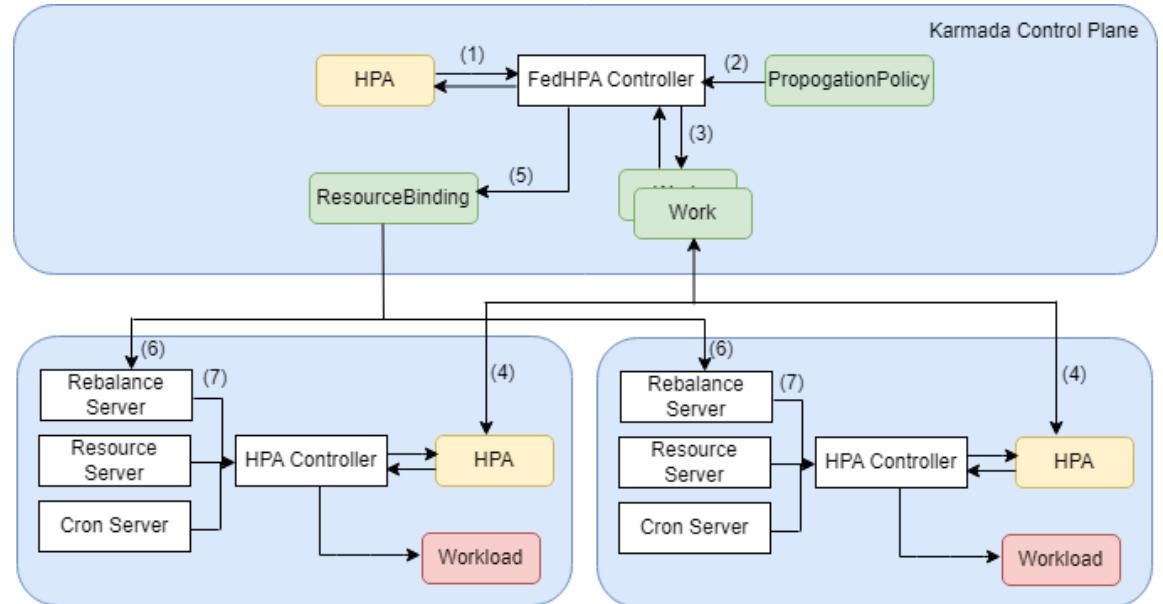
Architecture of HPA



- Components

- Processor
 - Resource Server
 - Cron Server
 - External Server
- Combinator
 - SmartHPAController(which is the name of HPA product in Trip.com, can be simply called HPAController)

Multi-cluster HPA based on Karmada



- Solution

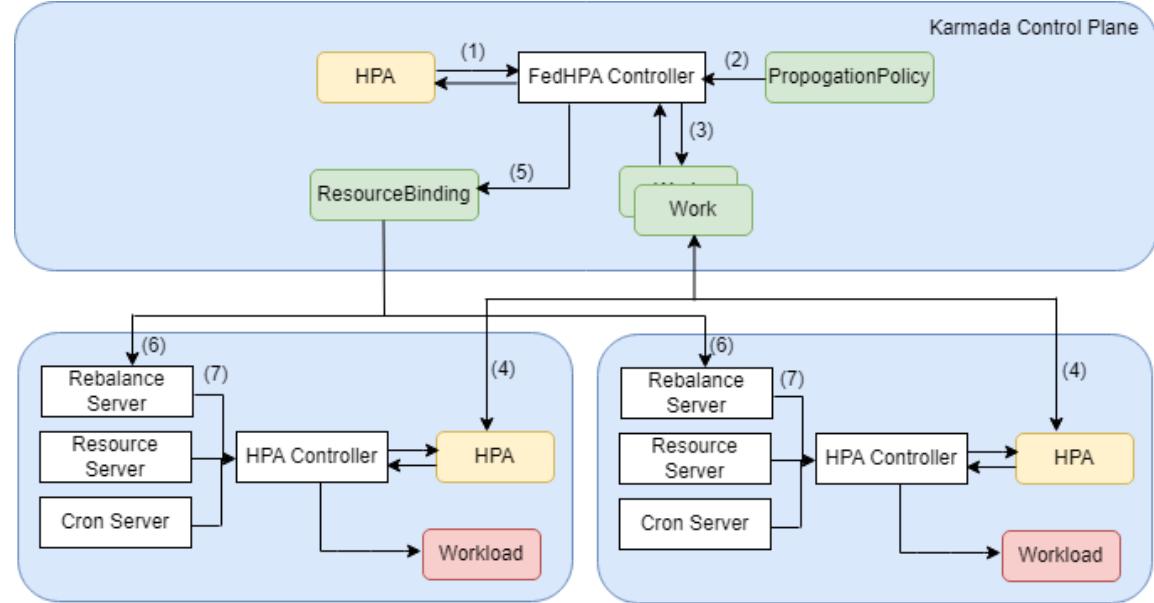
- The same HPA definition is shared in the federation and member cluster
- HPA in member clusters can work independently and will not be impacted due to the failure of other clusters or federation
- FedHPAController distributes HPA objects and rebalances between clusters according to the PropagationPolicy

design proposal: <https://github.com/karmada-io/karmada/blob/master/docs/proposals/hpa/federated-hpa.md>

```

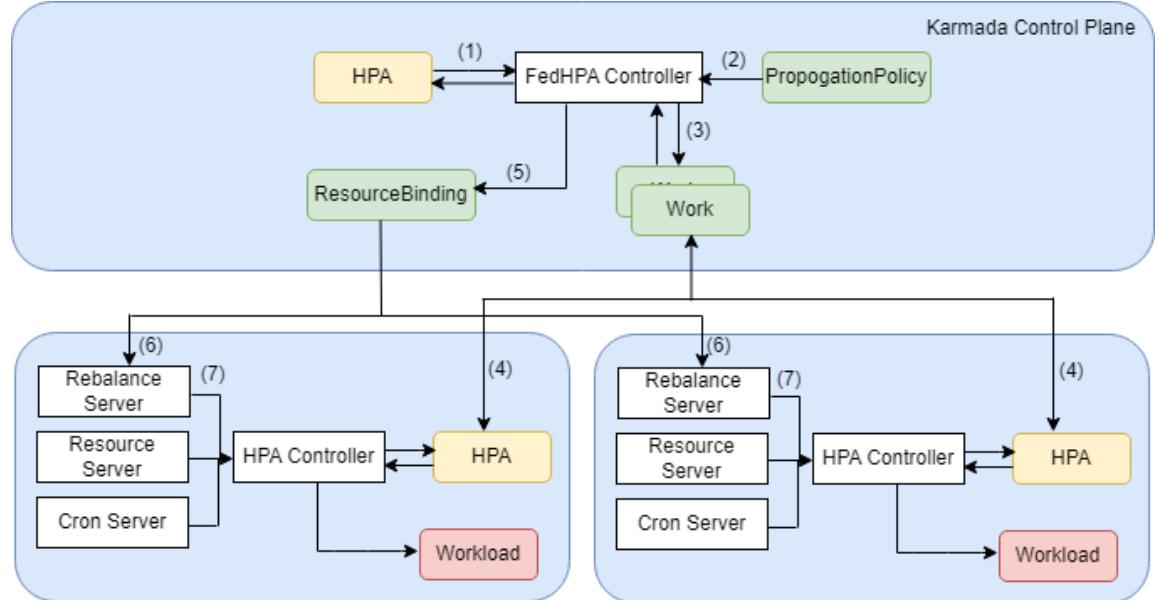
replicaScheduling:
  replicaDivisionPreference: Weighted
  replicaSchedulingType: Divided
  weightPreference:
    staticWeightList:
      - targetCluster:
          clusterNames:
            - shaxy-a
            weight: 3
      - targetCluster:
          clusterNames:
            - shaxy-b
            weight: 7
  priority: 0
  resourceSelectors:
    - apiVersion: captain.cloud.ctrip.com/v1
      kind: Rollout
      name: r100012822-91007693
      namespace: pro-captain
      schedulerName: default-scheduler
  
```

Distribute HPA Resource According to Propagation



- (1) FedHPAController watches HPA object
- (2) FedHPAController watches PropagationPolicy (weight)
- (3) FedHPAController creates Works according to the weight and HPA
 - work created for each HPA in each cluster
 - execute algorithm DistributeReplicasByWeight(replicas, clusterWeight)
- (4) Karmada distributes Works to member clusters

Rebalance between clusters gracefully



- How do the weights in Propagation affect the redistribution of minReplicas/maxReplicas among clusters?
 - FedHPAController calculates desired distribution according to the weight → (2)
- Reconcile multiple rounds to approach the desired weight
 - FedHPAController calculates the current round of applied distribution based on the actual ready pod and desired and updates the Works →(3)
 - FedHPAController distributes HPA to member clusters according to Works → (4)
 - FedHPAController calculates the current round of applied distribution based on the actual ready pod and desired and updates the Works →(3)
 -

Rebalance between clusters gracefully

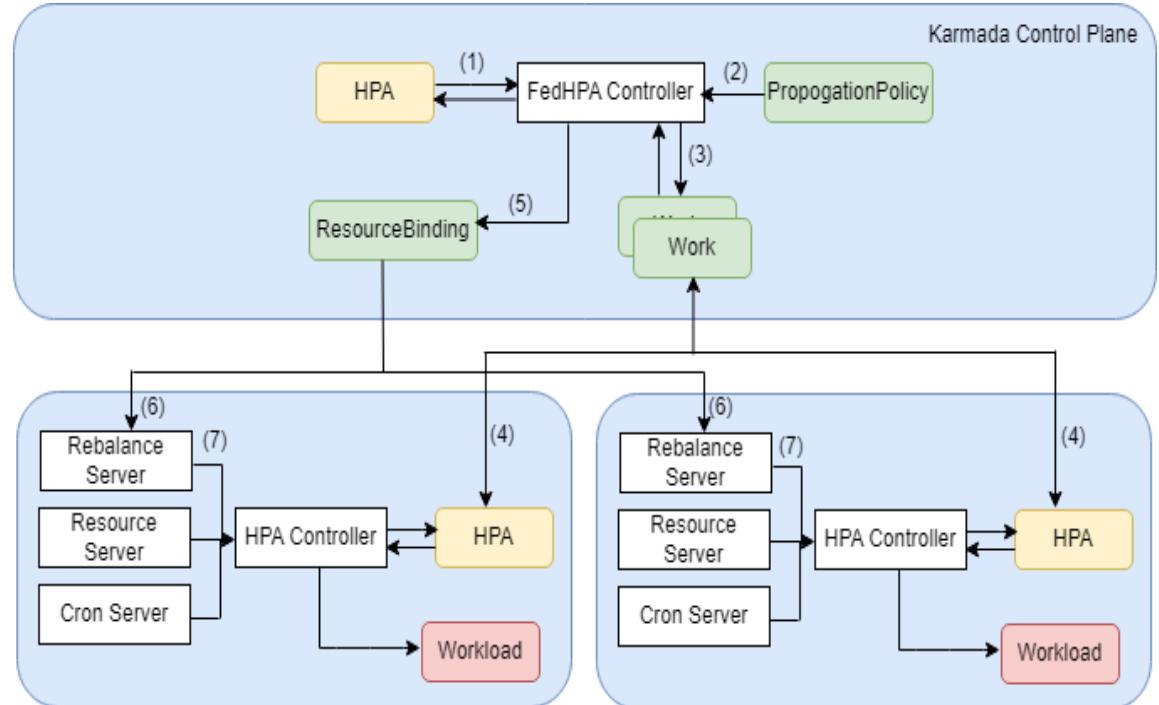
	before		=>	after	
	A	B		A	B
weight	1	0		0	1
desired max replicas	20	0		0	20
desired min replicas	10	0		0	10
actual replicas	10	0		0	10
ready replicas	10	0		0	0
note	-	-		10 in terminating	10 in creating

	Group X		
cluster	A	B	C
total minReplicas (0th remaining)	12		
weight	1	2	3
ready replicas (0th offset)	3	6	1
(1st) desired minReplicas	2	4	6
1st distribution	2	4	1
1st remaining	5		
1st offset (buffer / weight)	1	2	-
2nd desired	2	3	-
2nd distribution	1	2	-
2nd remaining	2		
2nd offset (buffer / weight)	-	-	5
3rd distribution	-	-	2
applied minReplicas	3	6	3

How to rebalance minReplicas/maxReplicas among member clusters?

- Calculate desired distribution for min/max replicas
- Reconcile multiple rounds to approach the desired distribution

Rebalance between clusters gracefully



- When the weight in the Propagation changes, how are the actual replicas redistributed among the clusters?
 - FedHPAController calculates the desired distribution based on actual replicas and weight → (2)
- Reconcile multiple rounds to approach the desired distribution
 - FedHPAController calculates the applied distribution based on actual and desired, recorded in the ResourceBinding annotation → (5)
 - RebanlaceServer within member cluster watches Resourcebinding, converts it into a rebalance signal → (6)
 - HPAController scales workload based on rebalance and other signals → (7)
 - FedHPAController calculates applied distribution and then saves it using annotation of ResourceBinding according to actual and desired values → (5)
 -

Rebalance between clusters gracefully

```

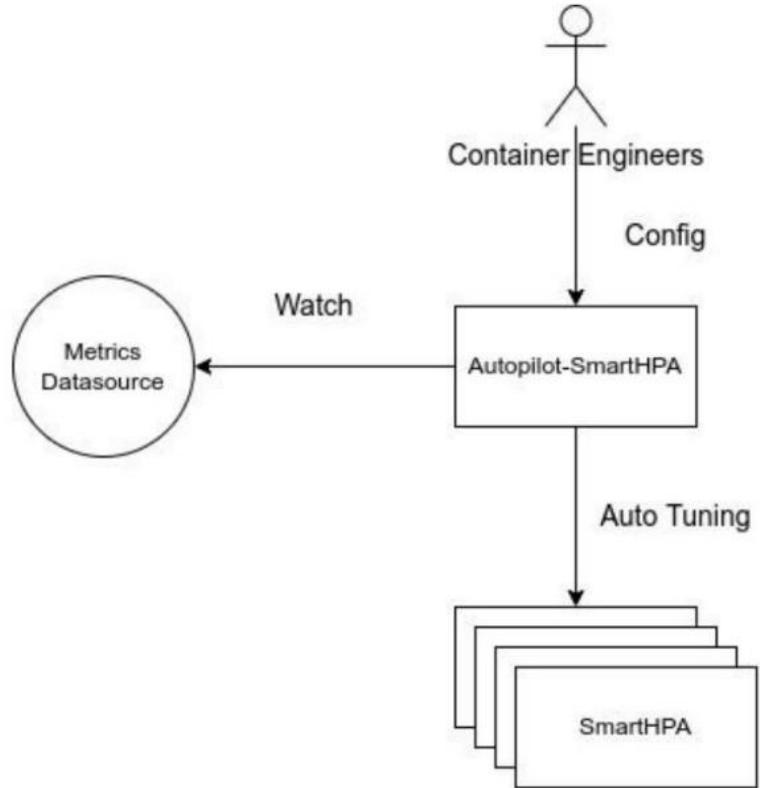
apiVersion: work.karmada.io/v1alpha2
kind: ResourceBinding
metadata:
  annotations:
    binding.karmada.io/customized-scheduling: enabled
    distribution.fed-smarthpa.karrier.io/desired-max-replicas: '{"shaxy-a":6,"shaxy-b":14}'
    distribution.fed-smarthpa.karrier.io/desired-min-replicas: '{"shaxy-a":1,"shaxy-b":2}'
    fed-rebalance-metric.fed-smarthpa.karrier.io/last-updated: "2023-09-25T11:33:47Z"
    fed-rebalance-metric.fed-smarthpa.karrier.io/metric-replicas: '{"shaxy-a":1,"shaxy-b":-1}'
    policy.karmada.io/applied-placement: '[{"clusterAffinity":{"clusterNames":["shaxy-a","shaxy-b"]}, "clusterTolerations":[{"key":"cluster.karmada.io/not-ready", "operator":"Exists", "effect":"NoExecute", "tolerationSeconds":300}, {"key":"cluster.karmada.io/unreachable", "operator":"Exists", "effect":"NoExecute", "tolerationSeconds":300}], "replicaScheduling":{"replicaSchedulingType":"Divided", "replicaDivisionPreference":"Weighted", "weightPreference":{"staticWeightList":[{"targetCluster":{"clusterNames":["shaxy-a"]}, "weight":3}, {"targetCluster":{"clusterNames":["shaxy-b"]}, "weight":7}]}}}'
  creationTimestamp: "2022-11-18T08:31:25Z"
  finalizers:
  - karmada.io/binding-controller
  generation: 21
  labels:
    binding.work.karmada.io/managed-by: fro-rb-controller
    propagationpolicy.karmada.io/name: app-r100012822-91007693
    propagationpolicy.karmada.io/namespace: pro-captain
  
```

How to balance the load between member clusters?

- calculate desired distribution according to actual replicas
- Reconcile multiple rounds to approach the desired distribution. Scale out the application on the side with fewer replicas than desired distribution first in each step.

cluster	A	B
weight	1	1
current replicas	9	1
desired replicas	5	5
rebalance signal	-1	3
current replicas	9	3
current replicas	7	3
rebalance signal	-1	4
current replicas	7	4
current replicas	6	4

How to manage HPA parameters?



- Fully managed parameters
 - min replicas
 - max replicas
 - target
 - ...
- Safety
 - support different profiles for different Apps
 - reconcile step by step and avoid big changes in a short time duration

Agenda

- Background
- Design
 - Reliability First
 - Capacity Scheduling
- Implementation
- Lessons learned

Lessons learned

- Performance and reliability are the cornerstones of elastic systems and deserve continued investment
- Establish data insight capabilities early
- Find influential seed users who share the same beliefs and solve difficult problems together, and other users will naturally follow
- The current downturn is temporary and we need to prepare in advance for future recovery



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2023

Thanks

Q&A